Big Data

Lecture 1. Introduction.

4th V does not only apply to Big Data (opinion) Same with 5th → added by business people to maintain hype.

big data ⇏ volume. velocity (streaming) ⇒ big data.

- Shared nothing: communication, consistency issues. (Spark)

Map/Reduce strategy → difficult!

- Lambda architecture: (λ): IoT example with trucks.

↳ divide data in batch (persistent; mutable) → master data it precomputes queries
speed (real time)
↳ serving layer.

Lecture 2. Intro to CS/Cloud.

Lecture 3. Dockerize model.

- Docker (SAAP): container management.
↳ has everything an app needs to run.

vs. Virtual Machines (good isolation // higher overload)

↓
slim
easy ship, pack, run
easy CI/CD
scability

- CLI: intermediario en las comunicaciones al servidor.

- docker: HOST

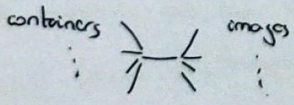

docker run -p (-v) 8080:8080 -v volume-path:mount jupyter/datascience-notebook
your disk inside container.

- Kubernetes: manages clusters of hosts (against docker images not having enough server)
↑ usable by anyone if app is dockerized.

Lecture 4. NoSQL : Not only SQL.

↳ scalable, fast, flexible, easy...
schemas (vs. SQL: rigid)

- Disadvantadge: no standarization rules (each DB has its own query language)
  bad with relational data
  ...

- Types: 1) Key value: big hash map (distributed) (key, value) pairs. (Amazon shopping basket)
  2) document-based: stores data as JSON objects. MongoDB. ↗ Wikipedia
  3) column-based: BigTable, Kassandra. SQL-likewise, difference in how it is stored. → we store rows → each one is HashMap (each can have its own columns)
  4) graph-based: model relationships between entities
     ↳ DATA-MODELING. generally not main DB.
     ↳ fraud detection.

Elastic Search → indexes to categorize similar data.
Hash-Map of Hash-Map.

SQL queries can crash with really large joins → MongoDB.
Check equivalencies Mongo - SQL.

Lecture 5. Hadoop & S3

- Hadoop Distributed File System (HDFS) + MapReduce (Computation)
  storage

do not your data to your local (huge!) ⇒ take your code to the server.

-blocks: storage units (files are divided in blocks).

data distribution
↳ if one goes down = corrupted data.

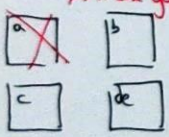

REPLICATE EACH BLOCK THRICE
↳ name node shows where each block is. (up to million of files)

- S3 : object storage (id, data, metadata) only to dump data there
  not splitted. unlimited.

use this for processing.

↳ Simple Storage Service

durable

two types of data: frequent or not accesses. → storage class.

- Standard
- Standard-IA, One Zone-IA (for backups)
- Intelligent Tiering (unknown, adaptative)
- Glacier: logs and archives.

- Hadoop Processing:
  - MapReduce: distribute by Keys after dividing into machines.
    (dividing into machines = map; distribute by Keys = reduce)

input splits → map → shuffle → reduce

Job/TaskTracker for status managements.

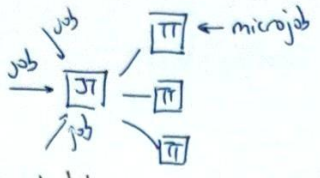

  - Hadoop 2: Yarn for resource management + job scheduling.

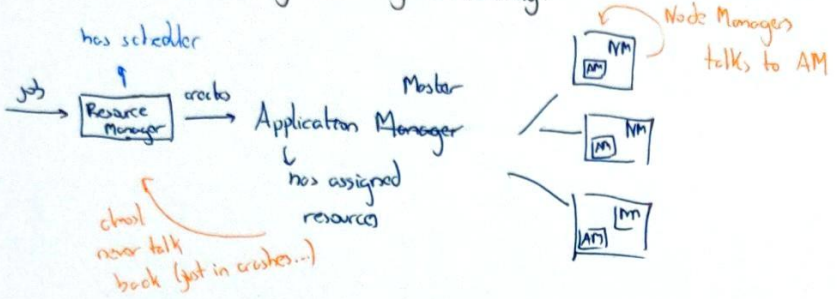

NM ≈ TT
AM ≈ JT (for one job only Hadoop)

Spark uses the same architecture.

16. Spark Fundamentals

- Not a replacement of Hadoop or BD storage: big data processing and analysis (instead of MapReduce)
  - Automatically distributes the work (fast)
  - Interactive exploration and analytics.
  - Easier code in more languages.

- Architecture:

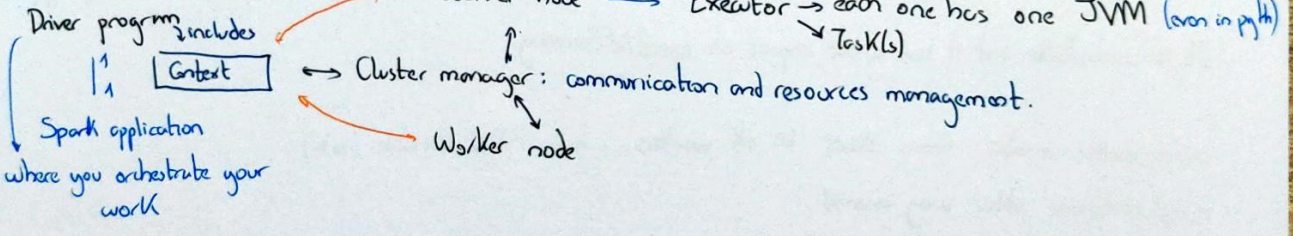

dbz: lazily = created only when needed.

- RDD: immutable collection that is partitioned and distributed among workers. → many types.
- Lineage: tree of transformations. → you can recreate a piece if lost and not all. > DAGs. no cycle graph.

transformations act on RDDs (data) / action: you want other kind of output count()
(lazy until you call an action)
(if you call twice an action, it will do everything again)

filter, coalesce (change number of packages)

(cost) ← save the RDD for reuse!
(cache())

- pair RDD: "dictionaries"

no loops in RDD (items are distributed)

reduce for operations with two items.

L8. Spark SQL

- Datasets and dataFrames
  - Typed: string? int? → typed.
  - untyped.
- SparkSession are the new entrypoint (instead of SparkContext)
  - ↳ better for SQL context.
  - DataFrames can be created from RDDs, hive, data sources
    - ↳ RDDs in the background.
- UDF in SQL: User defined functions.

L9. Spark Advanced

- Optimization: spark pipeline runs as much code as possible in a same stage.
  new stages have to be created when data needs to move across stages (shuffling)

Spark is immutable ⇒ each transformation ≡ a new RDD

- Narrow dependency (fast) each partition of parent RDD is used at most by one partition of child RDD

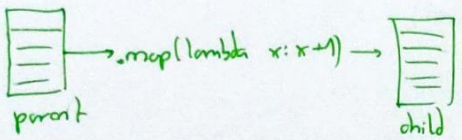

Wide " slow : " " " " " used ——— several " s of ...

It is controllable and it has great impact on speed/efficiency.

- broadcast variables: from driver to all workers. (non RDD variables (data))
- accumulators: other way around.

L10.

- Window SQL Function: group in a window a set of rows.
  - Spark SQL: window functions: 1) Ranking Functions
    2) Analytic Functions
    3)
- Streaming processing: spark streaming, Flink, apache storm.
  - ↳ microbatching: divide stream in batches. with DStreams. (Data Stream)

DStream

ou can access to specific RDDs:

we need