

Computer Vision

Lecture 1

• SIFT / HoG: Histogram of Gradients

Histogram of intensity does not work due to loss of geospatial information

↳ any permutation of original picture will work for the classifier.

Instead, consider a **Sobel Kernel** to detect horizontal and vertical edges → compute gradient in each position.

Then, plot a histogram based on how many times you perceive a specific value of angles between edge and grad.

• Gaussian vs. linear blur: Gaussian has more control through the value of σ .

• Filter banks: looks for particular lines and orientations. → to represent textures

Can be obtained through different derivatives of Gaussian filter.

• Mathematical morphology: (mostly works only on binary classification) → restoration, contour detection

structured element composed by 0's and 1's.

- 4 operations: 1) Dilation: if the center (white) interacts with another white pixel ⇒ apply logical sum (OR)
white sliding

2) Erosion: only when all whites in mask are matched with whites ⇒ Keep the center

3) Closing: dilate + erode. Equivalent to: if I can put the mask between two unconnected structures and it works as a bridge → connect them.

4) Opening: erode + dilate. Erode, but keep whole mask.

- Subtracting dilation and erosion ($Dilate(I) - Erosion(I)$) works as **contour detection**.

Lecture 2. Handcrafted Methods

define a descriptor → use machine learning


• Used when not enough data for Deep Learning.

• In some cases, descriptors can be injected into a NN.

- usable also in video.

Occlusion: common problem → but we still may have enough points of interests to detect on object!

To detect points of interest: compute a function that for each region, it computes an output for different scales.



→ the maximum represents stability for a scale → may be a PoI.
minimum also.

- **SIFT**: detects PoI and describes what is there.
very simple.

Feature vector with info about $grad_x$, $grad_y$, orientations, etc.

- robust towards global changes in illumination \rightarrow because it is normalized.

\downarrow it may find difficulties with local changes.

- not robust towards orientations: due to the axes of the gradients in the orientation.

\downarrow rotate the patches towards the predominant direction always.

(if math, they will be rotated in the same way \rightarrow orientation-friendly)

- robust towards affine transformations through **affine patches** (they use the two principal orthogonal directions)

useful when comparing in the same scene: cameras, robots, ...

- **RANSAC**: helps when there are too many candidate PoI with similar descriptions.

It will assume they will follow the same homography.

Check if all candidates follow the same transformation.



- **Highlight Features**: like what we used to detect faces



\rightarrow it helps us compute this 'easier'.

Integral images: cumulative sum from top-left.


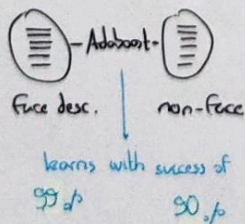
used in



= $\bullet - \bullet - \bullet + \bullet$
 $\downarrow \downarrow \downarrow \downarrow$
 in integral image (4 accesses)

- **Adaboost cascade**: - compute a lot of features from each image of faces

- we have



\downarrow next adaboost
 some faces + failed backgrounds

\rightarrow each classifier specializes in the error of the rest of classifiers

...
 success end of cascade:
 99.9% \sim 99.9%

- applied through sliding window:

- few operations (Integral images; most bgs don't go through many layers; only faces do)

Lecture 3. Convolutional Neural Network

↳ gradient descent to minimize the loss.

• **Stochastic Gradient Descent:** use sampling to optimize (use a batch instead of whole dataset)

↳ computational graph

• **Regularization:** way to guide towards no-overfitting.

We want weights to be regularized

$$\begin{pmatrix} 0.1 & 0.8 & 0.7 & \dots \\ 0.2 & 0.1 & 0.8 & \dots \\ 0.1 & & & \\ 0.5 & & & \\ \vdots & & & \end{pmatrix}$$

is better than

$$\begin{pmatrix} 0 & 1 & 1 & \dots \\ 0 & 0 & 1 & \dots \\ 0 & & & \\ 1 & & & \\ \vdots & & & \end{pmatrix}$$

↓
more robust toward changes in data.
Dropout has the same effect.

- small values of λ (please)

- **Data augmentation:** also in testing. crops, scales

• **Case studies:** 1) LeNet drawing in "Extra"

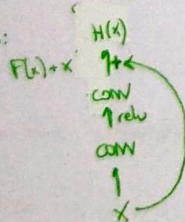
2) AlexNet: duplicated parallel structure (uncommon now)

3) VGG: 90M² params in FC. Simple structure

4) GoogleNet: "inception" module extracts different levels of detail
+ dimensionality reduction.

average pooling on output \Rightarrow no F.C layer at the end.

5) ResNet:



$$H(x) = F(x) + x$$

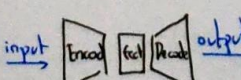
$F(x)$ is the residual

↑
shortcut for the grad \rightarrow no vanishing grads.

Lecture 5. Generative Models.

• Unsupervised learning \rightarrow understand the world.

• **Generative models:** generate samples from the same distribution. Many applications/interests.

• **Autoencoders:** do not use labels: 

You can change little values in the encoding to check what each component looks for \rightarrow interpretability.

- **GANs**: → it does not 'generalize' well. it is very dependant on the dataset distribution.

== == == == == == ==

- **Recurrent models:** based in previous states. generated by previous layers, generally.

with input of images: high res is way too unfeasible.
use representations (encoder, etc.)
↓
backbone

} all trained together.

- back propagation: keep a small window of the past.

in practice it is adding a new input to the formulation, multiplied by parameters (Wiv)

-attention: n-dimensional vector with activation of how important is each features.

-self-attention: multiply features * itself.

with attention: depending on which part of the captioning we are at, we may be interested in changing where I am looking at.

$$v = v^* \text{attention} \\ (v + \text{att_mask})$$

it comes from the previous stages.

Lecture 6. Human Pose Lecture.

- Decomposition in simpler parts \Rightarrow many false positives. (classical approach)
 \downarrow
 evaluate the joint probability of combinations.

- Human pose regression: $\text{img} \xrightarrow{\text{DL.}} 26 \text{ outputs (one per point)}$
 ↳ with bad regularization \rightarrow it will converge to the mean of the coordinates.

- cascade: the problem is partially solved \rightarrow use that as guide and start again
 dangers: overfitting
 converging to avg. again.

- Heatmaps: we want to regress an image with a gaussian centered on the joint. $\{ \rightarrow$ does not converge to mean.
- ↳ works infinitely better.

- no structure inference in the network. (Geometric Deep Learning) we need context (how to add?)
(\rightarrow convolutions?)

Context: * Stacked hourglass network: SoA, not efficient.

← inject.

- * Compositional models: optimize pairwise (between body parts) + bottom-top of hierarch. tree

• 3D: nothing really new \rightarrow refinement (3D shapes)

\hookrightarrow the projection in 2D can be generated from several 3D scenes (NP).

sharing problems but harder \rightarrow convergence to any again.

- multitask can be too complex

Just adapt what we do in 2D: (most suitable)

1) Geometric Heatmap

2) Predict 2D pose + add a component that learns 'z' coordinate.

== SoA ends here ==

• Problems we may find: occlusions and clothing variations. \Rightarrow data augmentation with occluders.

(paste random objects on the image. pixels ain't enough)

• Hand pose estimation (accessibility, gesture detection): do the same approach.

- PROBLEM: all parts ~~are~~ the same. You need hi-res and context.

\hookrightarrow look
 \hookrightarrow you do not really need RGB. It still works.

Voxel to voxel works best for 3D hand scenario.

• Bonus 1: human 3D \rightarrow clothing motion

• Bonus 2 - Face: affective computing \rightarrow what you feeling (games, neurorehabilitation,...)

It is not emotion recognition (we need neuron analysers)

detect faces commonly associated to...

\hookrightarrow Heatmap with 60+ joints in Face.

Lecture 7. Human behavior in video

• Basic approach: sliding window (+ time)

• Action / Gesture recognition: in general, 1 frame ain't enough.

\hookrightarrow For action understanding: gradients, bg extraction, optical flow

- Optical Flow: computed between 2 consecutive frames. It assumes pixels move slowly around its neighborhood.
pixels don't change 'much'
neighborhoods move in the same direction.

not very
uses.

$\left\{ \begin{array}{l} \text{equations} \\ \text{solver} \end{array} \right. \rightarrow$ shows translation of pixels.

"if something disappears" \rightarrow it should disappear smoothly.

The flow can be obtained in more efficient ways.

HoG + time \Rightarrow HoF (Flow)

against noise: "moving cubes" (like checking if hitbox moves)

- tracking: "following"

- STIPS: detect change in curvature to check changes in trajectories. Wonderful without noise.

- Levenshtein \rightarrow to compare trajectories.

adopt

length may be should not be penalized: invariant to time. (remove +1's in Levensh, but keep +diff)

- Dynamic Time Wrapping. divide in subproblems

- Deep Learning: using 3D kernels for convolutions \rightarrow one value. : if 3rd dimension is as big as temporal length
getting rid of temporal dim \rightarrow compute here.

\rightarrow 3D kernel output (spatial-temporal information)
3rd dim < time

5 strategies: 1) 2 stream ConvNets: one for image, one for optical flow.

2) 3D conv: use 3D kernels $\left\{ \begin{array}{l} \text{less params} \\ \text{than} \end{array} \right\} \left\{ \begin{array}{l} \text{more computations} \\ \text{than} \end{array} \right\} \rightarrow$ needs more training data than.

3) ConvNet + LSTM

4) Two-stream 3D convnets: 1) + 2) "if you lose temporal info, add 3D kernels" \rightarrow it may work.

5) Transformer-like: encoding spatio-temporal features.

1. Image processing

- Low level: images \rightarrow images
 - Mid level: images \rightarrow attributes
 - High level: scene understanding
- } img processing

- Color space (sRGB): strongly correlated, non-perceptual

- Image: $F(x,y) = z \rightarrow$ intensity / gray scale (digital if discrete)
- ↓
coord.

- digitalization: \equiv discretization
- resolution: number of pixels
- quantization: gray levels associated to b bits: $L = 2^b$
- brightness: global light intensity value $B = \frac{1}{N} \sum F(x,y)$
- contrast: perceived intensity difference between two regions $C = \frac{|f_a - f_b|}{f_a + f_b}$
- acutance: edge contrast
- sharpness: level of detail (acutance + resolution)
- noise: gaussian, salt & pepper

- Spatial filtering: operations in terms of neighbourhood (sliding window)
- } linearity ($F(a+b) = F(a) + F(b)$)
} shift invariant

box filter $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \rightarrow$ soothing

sharpening $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

sobel (edges): $\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

vertical horizontal

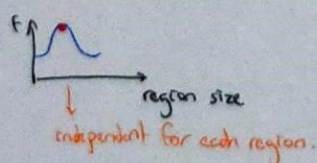
Gaussian / linear noise

- Segmentation: separate object from background
- 2D connectivity:
 - 4-conn: share a side
 - 8-conn: share a vertex

2. Handcrafted methods

- Detecting points of interest: bad feature if $\nabla^2 I$. Invariant to a lot of stuff.

- scale 'invariant' functions



$$F = \text{Kernel} * \text{image}$$

Laplacian

Diff of Gaussians (DIFT)

Affine patches (elliptical, two main eigenvec)

- Feature matching: describe PoIs.
 - Ransac
 - SIFT (128 dim: 6 orientations, 4x4 histogram array)

- Human detection:
 - HoG
 - Depth-based
 - Thermal: (intensity + gradient)
 - Fuse feature extraction of different modalities:
 - early \rightarrow concatenate
 - late \rightarrow weighted
 - middle \rightarrow include complementary feat. in middle steps of models.
- Iterative \neq fusion. Multi-modal.

L3. Introduction to CNN

• Linear classifier: $f(x; W) = Wx + b \rightarrow$ class scores \rightarrow use softmax (probability dist)

\downarrow image \downarrow squeezed \Rightarrow (n° pixels, 1)

- loss function: how good it currently is. (cross entropy)

- Neural Networks

1 layer: $F = Wx$

2 layer: $F = W_2 \max(0, W_1 x)$

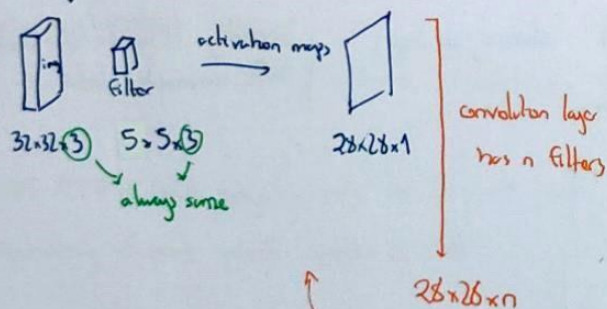
$\underbrace{\hspace{10em}}_{\text{non-linear}}$

$\begin{bmatrix} x & W_1 & n \end{bmatrix} \begin{bmatrix} 1 & 0 \end{bmatrix} \rightarrow \text{classes}$
 $(p(x, 1))$

• CNN:

img \rightarrow convolutions \rightarrow subsampling \rightarrow FCC \rightarrow output

stretched \times classes (as before)



stride: step taken in convolution's sliding window

zero padding: set borders to 0.

- ConvNet: sequence of convolutional layers: first layers \rightarrow low-level features (almost image-like)

mid level

high level

linearly separable classifier

applying non-linearity + pooling in middle

relu \downarrow activ. functions

sigmoid

tanh

ELU

maxout

(1x1 convs make sense)

- pooling: down sampling (max, avg) ... also sliding window

- weight init.

- batch normalization: after FC, or conv.

- learning rate: start with high and decay with time

- dropout

L4. Detection and segmentation.

- **Semantic segmentation**: label pixels with categories (ignore instances)

idea: use CNN maintaining size

or: 

pooling: near. neigh. $\begin{matrix} 1 & 2 \\ 3 & 4 \end{matrix} \rightarrow \begin{matrix} 11 & 22 \\ 33 & 44 \end{matrix}$

brd of ncls: $\begin{matrix} 1 & 2 & 10 & 20 \\ 3 & 4 & 30 & 40 \\ & & 00 & 00 \end{matrix}$

max pooling: remember position where max was.

transpose convolution.

- **Object detection**: deep learning applied to several crops of images (to detect of instances) (R-CNN)

selective search

- **Single Stage**: YOLO, SSD, Retina Net

image $\xrightarrow{\text{div.}}$ grid $\xrightarrow{\text{per cell}}$ 5 boxes $\xrightarrow{\text{progress}}$ (dx, dy, dh, dw, conf) $\xrightarrow{\text{predict}}$ class score $\xrightarrow{\text{by included}}$

- **Instance segmentation** Mask R-CNN

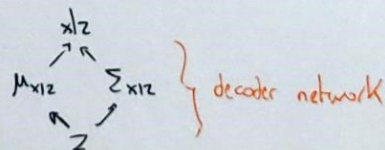
L5. Generative Models

- **Supervised**: learn a function $x \rightarrow y$
- **Unsupervised**: no labels \rightarrow discover structures (clustering, dim reduction...)

- **PixelRNN**: RNN (LSTM) that computes prob. dist of each pixel \downarrow

CNN: dependency of past pixels modeled by CNN

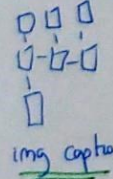
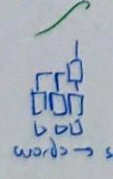
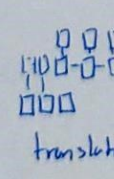

- **Variational Autoencoders**: probabilistic spin to Autoencoders



- **GAN**: loss \rightarrow train on $\frac{1}{\text{loss}}$

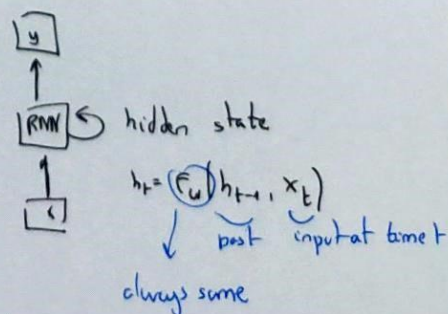
- can have convolutional architectures

L6. Recurrent Models

-    

- time can be long: truncate

idea



• LSTM: $\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} 0 \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$

$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

input, forget, output, gate

L7. Human pose lecture

L8. Human behaviour in video