

Constrained optimization: equality constraints

Lluís Garrido – lluis.garrido@ub.edu

November 2021

Abstract

This laboratory is focused on constrained optimization and, in particular, on equality constraints. The constrained optimization problem will be easily transformed into an unconstrained optimization problem using Lagrange multipliers.

1 Equality constraints: KKT conditions

Let us begin with a summary of equality constrained optimization¹.

Consider the problem of minimizing $f(\mathbf{x})$ subject to the constraints $h_i(\mathbf{x}) = 0$ for $i = 1 \dots m$, where $\mathbf{x} \in R^n$. Without any constraint, $m = 0$, the necessary condition for optimality is $\nabla f(\mathbf{x}) = 0$.

Let us now examine the case where $m = 1$, that is, a single constraint. With the constraint $h(\mathbf{x}) = 0$, we require that x lies on the graph of the (nonlinear) equation $h(\mathbf{x})$, see figure 1. Assume that \mathbf{x}^* is the optimal point we are looking for. The steepest descent direction at \mathbf{x}^* , $-\nabla f(\mathbf{x}^*)$, is orthogonal to the tangent of the contours of f through the point \mathbf{x}^* . A similar reasoning tells us that $\nabla h(\mathbf{x}^*)$ also is orthogonal to the curve $h(\mathbf{x}) = 0$ at \mathbf{x}^* . Otherwise, \mathbf{x}^* wouldn't be the optimum, see figure 1.

In addition, $\nabla f(\mathbf{x}^*)$ must be orthogonal to the tangent of the curve $h(\mathbf{x}) = 0$ at \mathbf{x}^* . This can be easily demonstrated and here we just give an insight. Assume that $c(t)$ is a curve $\{c = c(t) : t_0 \leq t \leq t_1\}$ such that $h(c(t)) = 0$. In other words, $c(t)$ is the feasible curve with respect to the constraint $h(\mathbf{x}) = 0$. For t^* , we have that $c(t^*) = \mathbf{x}^*$, the optimal value. Observe that $f(c(t))$ has a minimum at $t = t^*$, that is, the derivative of $f(c(t))$ with respect to t must vanish (i.e. be zero) at $t = t^*$. The derivative of $f(c(t))$ with respect to t at $t = t^*$ is

$$\left. \frac{d}{dt} f(c(t)) \right|_{t=t^*} = c'(t)^T \nabla f(c(t)) \Big|_{t=t^*} = c'(t^*)^T \nabla f(\mathbf{x}^*)$$

The previous expression indicates us that $\nabla f(\mathbf{x}^*)$ must be orthogonal to the tangent of the curve $h(\mathbf{x}) = 0$ at $\mathbf{x} = \mathbf{x}^*$.

Therefore, $\nabla f(\mathbf{x}^*)$ and $\nabla h(\mathbf{x}^*)$ must both lie along the same line. That is, for some $\lambda \in R$ we must have

$$\nabla f(\mathbf{x}^*) = \lambda \nabla h(\mathbf{x}^*)$$

Thus, if \mathbf{x}^* is a minimizer, the necessary condition reduces to

$$\nabla f(\mathbf{x}^*) - \lambda \nabla h(\mathbf{x}^*) = 0$$

¹The text written in this section has been obtained from Griva, I.; Nash, S.; Sofer, A., "Linear and nonlinear optimization", SIAM and <http://www.pitt.edu/~jrc/class/opt/notes3.pdf>.




Figure 1: For a single constraint, $m = 1$, we seek a point \mathbf{x}^* such that $\nabla f(\mathbf{x}^*) = \lambda \nabla h(\mathbf{x}^*)$.

For the general case, in which we have the equality constraints $h_i(\mathbf{x}) = 0$ for $j = 1 \dots m$, the above necessary condition should hold for each constraint. Assuming that $\nabla h_i(\mathbf{x}^*)$ are linearly independent (or equivalently, the Jacobian matrix has full row rank), the point $\mathbf{x}^* \in R^n$ must satisfy the following necessary condition for some $\lambda^* \in R^m$, that is

$$\nabla f(\mathbf{x}^*) - \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*) = 0 \quad h_i(\mathbf{x}^*) = 0 \quad \forall i$$

The latter conditions are known as the **Karush-Kuhn-Tucker** (KKT) conditions. The previous necessary condition can be easily extended to inequalities.

The Lagrangian is defined to be

$$\mathcal{L}(\mathbf{x}, \lambda) = f(\mathbf{x}) - \sum_{i=1}^m \lambda_i h_i(\mathbf{x})$$

In order to find the points that satisfy the equality conditions we may think that a gradient descent can be applied to the Lagrangian function. Unfortunately, the optimum value \mathbf{x}^* is in general a **saddle point** of the Lagrangian – not a minimum of $\mathcal{L}(\mathbf{x}, \lambda)$ – as the following example shows.

Example

Consider the one dimensional problem of minimizing $f(x) = x^2$ subject to the constraint $x - 1 = 0$. The Lagrangian function is

$$\mathcal{L}(x, \lambda) = x^2 - \lambda(x - 1)$$

The solution to this problem is $x^* = 1$, with Lagrange multiplier $\lambda^* = 2$. Since $\nabla \mathcal{L}(x, \lambda) = (2x - \lambda, -(x - 1))$, then $\nabla \mathcal{L}(1, 2) = (0, 0)^T$, and indeed (x^*, λ^*) is a stationary point of the Lagrangian. The Hessian matrix of the Lagrangian is

$$\nabla^2 \mathcal{L}(x^*, \lambda^*) = \begin{pmatrix} 2 & -1 \\ -1 & 0 \end{pmatrix}$$




Figure 2: Lagrangian function, see example. This image has been obtained from I. Griva, S.G. Nash, A. Sofer, “Linear and Nonlinear Optimization”.

This is an indefinite matrix. Thus (x^*, λ^*) is a saddle point of the Lagrangian function, as can be seen in figure 2.

There are two approaches that can be used to ensure convergence. The first approach attempts to ensure that \mathbf{x}^k is feasible at each iteration (i.e. the constraints are satisfied at each iteration). This is the approach that is going to be taken in this lab. The second approach constructs a new function, related to the Lagrangian, that (ideally) has a minimum at $(\mathbf{x}^*, \lambda^*)$. Whereas the former are called feasible-point methods, the second are known as penalty and barrier methods.

2 Sequential Quadratic Programming

Sequential quadratic programming is a popular and successful technique for solving nonlinearly constrained problems. The main idea is to obtain, for each iteration, a search direction by solving an approximation of the original problem. In particular, for each iteration we will solve a quadratic program, that is, a problem with a quadratic objective function and linear constraints.

Our focus is to solve (we focus only on one constraint)

$$\begin{aligned} &\text{minimize} && f(\mathbf{x}) \\ &\text{subject to} && h(\mathbf{x}) = 0 \end{aligned}$$

The Lagrangian is

$$\mathcal{L}(\mathbf{x}, \lambda) = f(\mathbf{x}) - \lambda h(\mathbf{x})$$

The first order optimality condition is

$$\nabla \mathcal{L}(\mathbf{x}, \lambda) = 0$$

Assume that we use the Newton method to minimize the Lagrange function

$$\begin{pmatrix} \mathbf{x}^{k+1} \\ \lambda^{k+1} \end{pmatrix} = \begin{pmatrix} \mathbf{x}^k \\ \lambda^k \end{pmatrix} + \alpha^k \begin{pmatrix} \mathbf{d}^k \\ \nu^k \end{pmatrix} \tag{1}$$

where \mathbf{d}^k and ν^k are obtained as the solution to the Newton linear system

$$\nabla^2 \mathcal{L}(\mathbf{x}^k, \lambda^k) \begin{pmatrix} \mathbf{d}^k \\ \nu^k \end{pmatrix} = -\nabla \mathcal{L}(\mathbf{x}^k, \lambda^k)$$

This linear system has the form

$$\begin{pmatrix} \nabla_{xx}^2 \mathcal{L}(\mathbf{x}^k, \lambda^k) & -\nabla h(\mathbf{x}^k) \\ -\nabla h(\mathbf{x}^k)^T & 0 \end{pmatrix} \begin{pmatrix} \mathbf{d}^k \\ \nu^k \end{pmatrix} = \begin{pmatrix} -\nabla_x \mathcal{L}(\mathbf{x}^k, \lambda^k) \\ h(\mathbf{x}^k) \end{pmatrix} \quad (2)$$

The latter system of equations represents the first-order optimality conditions of the following optimization problem

$$\begin{aligned} \text{minimize} \quad & q(\mathbf{d}) = \frac{1}{2} \mathbf{d}^T \left[\nabla_{xx}^2 \mathcal{L}(\mathbf{x}^k, \lambda^k) \right] \mathbf{d} + \mathbf{d}^T \left[\nabla_x \mathcal{L}(\mathbf{x}^k, \lambda^k) \right] \\ \text{subject to} \quad & \left[\nabla h(\mathbf{x}^k)^T \right] \mathbf{d} + h(\mathbf{x}^k) = 0 \end{aligned} \quad (3)$$

This optimization problem is a quadratic program; that is, it is the minimization of a quadratic function subject to linear constraints. The quadratic function is a Taylor series approximation to the Lagrangian at $(\mathbf{x}^k, \lambda^k)$ and the constraints are a linear approximation to $h(\mathbf{x}^k + \mathbf{d}) = 0$.

In a Sequential Quadratic Optimization method, at each iteration the linear system of equation (2) is solved to obtain (\mathbf{d}^k, ν^k) . These are used to update $(\mathbf{x}^k, \lambda^k)$, and the process repeats at the new point.

Example

We apply the Sequential Quadratic Optimization method to the problem

$$\begin{aligned} \text{minimize} \quad & f(x_1, x_2) = e^{3x_1} + e^{-4x_2} \\ \text{subject to} \quad & h(x_1, x_2) = x_1^2 + x_2^2 - 1 = 0 \end{aligned}$$

The solution to this problem is $(x_1^*, x_2^*) \approx (-0.74834, 0.66332)^T$ with $\lambda^* \approx -0.21233$, which in fact is a **saddle point**! The Hessian matrix of the Lagrangian has two positive and one negative eigenvalues.

Let us see how to proceed with the first iteration with the initial guess $\mathbf{x}^0 = (x_1^0, x_2^0) = (-1, 1)^T$ and $\lambda^0 = -1$. At this point the first order derivatives are

$$\begin{aligned} \nabla f &= \begin{pmatrix} 0.14936 \\ -0.07326 \end{pmatrix} \\ \nabla h &= \begin{pmatrix} -2 \\ 2 \end{pmatrix} \\ \nabla_x \mathcal{L} &= \nabla f - \lambda \nabla h = \begin{pmatrix} -1.85064 \\ 1.92674 \end{pmatrix} \end{aligned}$$

and the second order derivatives are

$$\nabla^2 f = \begin{pmatrix} 0.44808 & 0 \\ 0 & 0.29305 \end{pmatrix}$$

$$\begin{aligned}\nabla^2 h &= \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix} \\ \nabla_{xx}^2 \mathcal{L} &= \nabla^2 f - \lambda \nabla^2 h = \begin{pmatrix} 2.44808 & 0 \\ 0 & 2.29305 \end{pmatrix}\end{aligned}$$

The corresponding quadratic program is given by Eq. (3) and its solution can be found using the optimality conditions given by Eq. (2)

$$\begin{pmatrix} 2.44808 & 0 & 2 \\ 0 & 2.29305 & -2 \\ 2 & -2 & 0 \end{pmatrix} \begin{pmatrix} d_1^0 \\ d_2^0 \\ \nu^0 \end{pmatrix} = \begin{pmatrix} 1.85064 \\ -1.92674 \\ 1 \end{pmatrix}$$

The solution of the quadratic program is

$$\begin{aligned}\mathbf{d}^0 &= \begin{pmatrix} d_1^0 \\ d_2^0 \end{pmatrix} = \begin{pmatrix} 0.22577 \\ -0.27423 \end{pmatrix} \\ \nu^0 &= (0.64896)\end{aligned}$$

We may then perform the update of the parameters $(\mathbf{x}^k, \lambda^k)$ using Eq. (1) with a step $\alpha = 1$. The new estimate of the solution is

$$\begin{aligned}\mathbf{x}^1 &= \mathbf{x}^0 + \mathbf{d}^0 = \begin{pmatrix} -0.77423 \\ 0.72577 \end{pmatrix} \\ \lambda^1 &= \lambda^0 + \nu^0 = (-0.35104)\end{aligned}$$

This procedure is repeated until the problem converges to the solution.

Proposed experiments

1. One simple way to proceed is to take $\alpha^k = 1$ and iteratively update the current point to obtain the next. This is a simple way to proceed that is proposed to perform first. The stopping condition should be performed over $\nabla_x \mathcal{L}$. Test this approach and check if it works using the starting point proposed in the example.
2. This basic iteration also has drawbacks, leading to a number of vital questions. It is a Newton-like iteration, and thus may diverge from poor starting points. In our example we have started from a point that is near to the optimal solution. Try to perform some experiments with starting points that are farther away of the optimal solution.
3. One way to find the optimal solution from points that are far away of the optimal solution is to start the optimization with another function that allows us to find an approximation to the solution we are looking for. Once an approximate solution is found, we can apply the Newton-based technique we have presented previously to find the optimal solution.

The function that allows us to find an approximation to the solution we are looking for is called, in this context, the merit function. Usually, a merit function is the sum of terms that include the objective function and the amount of infeasibility of the constraints. One

example of a merit function for the problem we are treating is the quadratic penalty function (i.e. constraints are penalized quadratically)

$$\mathcal{M}(x_1, x_2) = f(x_1, x_2) + \rho h(x_1, x_2)^2$$

where ρ is some positive number. The greater the value of ρ , the greater the penalty for infeasibility. The difficulty arises in defining a proper merit function for a particular equality constrained problem.

Here we propose you to take $\rho = 10$ (thus, we penalize a lot the constraint) and perform a classical gradient descent (with backtracking if you want) to find an approximation to the solution we are looking for. Observe if you arrive near to the optimal solution of the problem. Take into account that you may have numerical problems with the gradient. A simple way to deal with it is to normalize the gradient at each iteration, $\nabla \mathcal{M}(x) / \|\nabla \mathcal{M}(x)\|$, and use this normalized gradient as search direction.

4. As previously commented, the minimizers of the merit function $\mathcal{M}(x_1, x_2)$ do not necessarily have to coincide with the minimizers of the constrained problem. Thus, once we “sufficiently” approach the optimal solution we may use the Newton method (with $\alpha = 1$) to find the solution to the problem.

Therefore the algorithm consists in starting with the Merit function to obtain an approximation to the optimal point we are looking for. Once an approximation to the solution is found, use the Newton-based method to find the optimal solution. Check if you are able to find the optimal solution to the problem.

Report

You are asked to deliver a report (PDF, notebook, or whatever else you prefer) that can be done in *pairs of two persons* (or *individually*). Comment each of the steps you have followed as well as the results and plots you obtain. Do not expect the reader (i.e. me) to interpret the results for you. I would like to see if you are able to understand the results you have obtained.

If you want to include some parts of code, please include it within the report. Do not include it as separate files. You may just deliver the Python notebook if you want.